

In [1]:

```
import warnings
warnings.filterwarnings("ignore")
from sklearn.datasets import load_boston
from random import seed
from random import randrange
from csv import reader
from math import sqrt
from sklearn import preprocessing
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from prettytable import PrettyTable
from sklearn.linear_model import SGDRegressor
from sklearn import preprocessing
from sklearn.metrics import mean_squared_error, mean_absolute_error
from sklearn.model_selection import train_test_split
```

In [2]:

```
# Citation
# https://www.kaggle.com/premvardhan/stochasticgradientdescent-implementation-lr-python
# https://www.kaggle.com/arpandas65/simple-sgd-implementation-of-linear-regression
```

In [3]:

```
boston = load_boston()
print(boston.DESCR)
```

.. _boston_dataset:

Boston house prices dataset

****Data Set Characteristics:****

:Number of Instances: 506

:Number of Attributes: 13 numeric/categorical predictive. Median Value (attribute 14) is usually the target.

:Attribute Information (in order):

- CRIM per capita crime rate by town
- ZN proportion of residential land zoned for lots over 25,000 sq.ft.
- INDUS proportion of non-retail business acres per town
- CHAS Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)
- NOX nitric oxides concentration (parts per 10 million)
- RM average number of rooms per dwelling
- AGE proportion of owner-occupied units built prior to 1940
- DIS weighted distances to five Boston employment centres
- RAD index of accessibility to radial highways
- TAX full-value property-tax rate per \$10,000
- PTRATIO pupil-teacher ratio by town
- B 1000(Bk - 0.63)^2 where Bk is the proportion of blacks by town
- LSTAT % lower status of the population
- MEDV Median value of owner-occupied homes in \$1000's

:Missing Attribute Values: None

:Creator: Harrison, D. and Rubinfeld, D.L.

This is a copy of UCI ML housing dataset.

<https://archive.ics.uci.edu/ml/machine-learning-databases/housing/>

This dataset was taken from the StatLib library which is maintained at Carnegie Mellon University.

The Boston house-price data of Harrison, D. and Rubinfeld, D.L. 'Hedonic prices and the demand for clean air', J. Environ. Economics & Management, vol. 5, 81-102, 1978. Used in: Belsley, Kuh & Welsch, 'Regression diagnostics

vol.3, 61-102, 1976. Used in Belsley, Kuh & Welsch, 'Regression diagnostics ...', Wiley, 1980. N.B. Various transformations are used in the table on pages 244-261 of the latter.

The Boston house-price data has been used in many machine learning papers that address regression problems.

.. topic:: References

- Belsley, Kuh & Welsch, 'Regression diagnostics: Identifying Influential Data and Sources of Collinearity', Wiley, 1980. 244-261.
- Quinlan, R. (1993). Combining Instance-Based and Model-Based Learning. In Proceedings on the Tenth International Conference of Machine Learning, 236-243, University of Massachusetts, Amherst. Morgan Kaufmann.

In [4]:

```
# Split train and test data
X = boston.data
Y = boston.target

X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.33, random_state=5)
```

In [5]:

```
# Create Data Frame
boston_data = pd.DataFrame(boston.data, columns=boston.feature_names)

# show data
boston_data.head()
```

Out[5]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33

In [6]:

```
boston_data.shape
```

Out[6]:

(506, 13)

In [7]:

```
# Standrize data
scaler = preprocessing.StandardScaler()

scaler.fit(X_train)

X_train = scaler.transform(X_train)
X_test = scaler.transform(X_test)
```

In [8]:

```
print(X_train.shape)
print(X_test.shape)
print(Y_train.shape)
print(Y_test.shape)
```

(339, 13)
(167, 13)

```
(10/, 15)
(339,)
(167,)
```

In [9]:

```
# Implement Linear Regression
# code source: https://medium.com/@haydar\_ai/learning-data-science-day-9-linear-regression-on-boston-housing-dataset-cd62a80775ef
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
import matplotlib.pyplot as plt

lm = LinearRegression()
lm.fit(X_train, Y_train)

Y_pred = lm.predict(X_test)

print('Coefficients: \n', lm.coef_)
# The mean squared error
print("Mean squared error: %.2f" % mean_squared_error(Y_test, Y_pred))
# Explained variance score: 1 is perfect prediction
# print("R^2 score: %.2f" % lm.score(X_test, Y_test))
print('Variance score: %.2f' % r2_score(Y_test, Y_pred))

plt.scatter(Y_test, Y_pred)
plt.xlabel("Prices: $Y_i$")
plt.ylabel("Predicted prices: $\hat{Y}_i$")
plt.title("Prices vs Predicted prices: $Y_i$ vs $\hat{Y}_i$")
plt.show()
```

Coefficients:
[-1.31193031 0.86187745 -0.16719287 0.18957843 -1.48658584 2.79131565
-0.32737703 -2.77204093 2.97567549 -2.2727549 -2.13375869 1.05842993
-3.33495407]
Mean squared error: 28.53
Variance score: 0.70




In [10]:

```
# First implement SKLearn SGD
# SkLearn SGD classifier
clf_ = SGDRegressor()
clf_.fit(X_train, Y_train)

plt.scatter(Y_test, clf_.predict(X_test))
plt.grid()
plt.xlabel('Actual y')
plt.ylabel('Predicted y')
plt.title('scatter plot between actual y and predicted y')
plt.show()
print('Mean Squared Error :', mean_squared_error(Y_test, clf_.predict(X_test)))
print('Mean Absolute Error :', mean_absolute_error(Y_test, clf_.predict(X_test)))
```

scatter plot between actual y and predicted y



Mean Squared Error : 28.545951396225185
Mean Absolute Error : 3.450685306695088

In [11]:

```
# SkLearn SGD classifier predicted weight matrix
sklearn_w=clf_.coef_
sklearn_w
```

Out[11]:

```
array([-1.19998397,  0.70823758, -0.47154859,  0.24990282, -1.25091027,
        2.89665502, -0.38369813, -2.64985472,  1.78730939, -1.05960085,
       -2.08913235,  1.03175028, -3.31171865])
```

In [12]:

```
# Implement Custom SGD
boston_data.describe()
```

Out[12]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTR
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634	68.574901	3.795043	9.549407	408.237154	18.461771
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617	28.148861	2.105710	8.707259	168.537116	2.107673
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	2.900000	1.129600	1.000000	187.000000	12.600000
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500	45.025000	2.100175	4.000000	279.000000	17.400000
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500	77.500000	3.207450	5.000000	330.000000	19.000000
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500	94.075000	5.188425	24.000000	666.000000	20.200000
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	100.000000	12.126500	24.000000	711.000000	22.000000

In [13]:

```
boston_data = ((boston_data - boston_data.mean())/boston_data.std())
boston_data.head()
```

Out[13]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
0	0.419367	0.284548	1.286636	0.272329	0.144075	0.413263	0.119895	0.140075	0.981871	0.665949	1.457558	0.440616	1.074499
1	0.416927	0.487240	0.592794	0.272329	0.739530	0.194082	0.366803	0.556609	0.867024	0.986353	0.302794	0.440616	0.491953
2	0.416929	0.487240	0.592794	0.272329	0.739530	1.281446	0.265549	0.556609	0.867024	0.986353	0.302794	0.396035	1.207532
3	0.416338	0.487240	1.305586	0.272329	0.834458	1.015298	0.809088	1.076671	0.752178	1.105022	0.112920	0.415751	1.360171

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
4	0.412074	0.487240	1.305586	0.272329	0.834458	1.227362	0.510674	1.076671	0.752178	1.105022	0.112920	0.440616	1.025487

In [14]:

```
boston_data["PRICE"] = boston.target

boston_data.head()
# X = boston_data
# Y = boston.target

# print(X.shape)
# print(Y.shape)

# X.head()
# Y.head()
```

Out[14]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
0	0.419367	0.284548	1.286636	0.272329	0.144075	0.413263	0.119895	0.140075	0.981871	0.665949	1.457558	0.440616	1.074499
1	0.416927	0.487240	0.592794	0.272329	0.739530	0.194082	0.366803	0.556609	0.867024	0.986353	0.302794	0.440616	0.491953
2	0.416929	0.487240	0.592794	0.272329	0.739530	1.281446	0.265549	0.556609	0.867024	0.986353	0.302794	0.396035	1.207532
3	0.416338	0.487240	1.305586	0.272329	0.834458	1.015298	0.809088	1.076671	0.752178	1.105022	0.112920	0.415751	1.360171
4	0.412074	0.487240	1.305586	0.272329	0.834458	1.227362	0.510674	1.076671	0.752178	1.105022	0.112920	0.440616	1.025487

In [15]:

```
Y = boston_data['PRICE']
X = boston_data.drop('PRICE', axis=1)

print(X.shape)
print(Y.shape)

print(X.head())
print("*"*100)
print(Y.head())
```

```
(506, 13)
(506,)

      CRIM      ZN      INDUS      CHAS      NOX      RM      AGE \
0 -0.419367  0.284548 -1.286636 -0.272329 -0.144075  0.413263 -0.119895
1 -0.416927 -0.487240 -0.592794 -0.272329 -0.739530  0.194082  0.366803
2 -0.416929 -0.487240 -0.592794 -0.272329 -0.739530  1.281446 -0.265549
3 -0.416338 -0.487240 -1.305586 -0.272329 -0.834458  1.015298 -0.809088
4 -0.412074 -0.487240 -1.305586 -0.272329 -0.834458  1.227362 -0.510674

      DIS      RAD      TAX      PTRATIO      B      LSTAT
0  0.140075 -0.981871 -0.665949 -1.457558  0.440616 -1.074499
1  0.556609 -0.867024 -0.986353 -0.302794  0.440616 -0.491953
2  0.556609 -0.867024 -0.986353 -0.302794  0.396035 -1.207532
3  1.076671 -0.752178 -1.105022  0.112920  0.415751 -1.360171
4  1.076671 -0.752178 -1.105022  0.112920  0.440616 -1.025487
*****

0    24.0
1    21.6
2    34.7
3    33.4
4    36.2
Name: PRICE, dtype: float64
```

In [16]:

```
x_train, x_test, y_train, y_test = train_test_split(X, Y, test_size = 0.3)
```

```
print(x_train.shape, x_test.shape, y_train.shape, y_test.shape)
```

```
(354, 13) (152, 13) (354,) (152,)
```

In [17]:

```
x_train["PRICE"] = y_train
```

In [18]:

```
def cost_function(b, m, features, target):
    totalError = 0
    for i in range(0, len(features)):
        x = features
        y = target
        #error = (y-ypred)^2
        totalError += (y[:,i] - (np.dot(x[i] , m) + b)) ** 2
    return totalError / len(x)
```

In [19]:

```
def absolute_cost_function(b, m, features, target):
    totalError = 0
    for i in range(0, len(features)):
        x = features
        y = target
        #error = (y-ypred)
        totalError += (y[:,i] - (np.dot(x[i] , m) + b))
    return totalError / len(x)
```

In [20]:

```
# The total sum of squares (proportional to the variance of the data) i.e. ss_tot
# The sum of squares of residuals, also called the residual sum of squares i.e. ss_res
# the coefficient of determination i.e. r^2 (r squared)
def r_sq_score(b, m, features, target):
    for i in range(0, len(features)):
        x = features
        y = target
        mean_y = np.mean(y)
        ss_tot = sum((y[:,i] - mean_y) ** 2)
        # sum((y-(mx+b)))^2
        ss_res = sum(((y[:,i]) - (np.dot(x[i], m) + b)) ** 2)
        r2 = 1 - (ss_res / ss_tot)
    return r2
```

In [21]:

```
def gradient_decent(w0, b0, train_data, x_test, y_test, learning_rate):
    n_iter = 500
    partial_deriv_m = 0
    partial_deriv_b = 0
    cost_train = []
    cost_test = []
    for j in range(1, n_iter):

        # Train sample
        train_sample = train_data.sample(160)
        y = np.asmatrix(train_sample["PRICE"])
        x = np.asmatrix(train_sample.drop("PRICE", axis = 1))
        # Test sample
        #x_test["PRICE"] = [y_test]
        #test_data = x_test
        #test_sample = test_data.sample()
        #y_test = np.asmatrix(test_sample["PRICE"])
        #x_test = np.asmatrix(test_sample.drop("PRICE", axis = 1))

        for i in range(len(x)):
            partial_deriv_m += np.dot(-2*x[i].T , (y[:,i] - np.dot(x[i] , w0) + b0))
            partial_deriv_b += -2*(y[:,i] - (np.dot(x[i] , w0) + b0))
```

```

w1 = w0 - learning_rate * partial_deriv_m
b1 = b0 - learning_rate * partial_deriv_b

if (w0==w1).all():
    #print("W0 are\n", w0)
    #print("\nW1 are\n", w1)
    #print("\n X are\n", x)
    #print("\n y are\n", y)
    break
else:
    w0 = w1
    b0 = b1
    learning_rate = learning_rate/2

error_train = cost_function(b0, w0, x, y)
cost_train.append(error_train)
error_test = cost_function(b0, w0, np.asmatrix(x_test), np.asmatrix(y_test))
cost_test.append(error_test)

#print("After {0} iteration error = {1}".format(j, error_train))
#print("After {0} iteration error = {1}".format(j, error_test))

return w0, b0, cost_train, cost_test

```

In [22]:

```

# Run our model
learning_rate = 0.001
w0_random = np.random.rand(13)
w0 = np.asmatrix(w0_random).T
b0 = np.random.rand()

optimal_w, optimal_b, cost_train, cost_test = gradient_decent(w0, b0, x_train, x_test, y_test, learning_rate)
print("Coefficient: {} \n y_intercept: {}".format(optimal_w, optimal_b))

'''
error = cost_function(optimal_b, optimal_w, np.asmatrix(x_test), np.asmatrix(y_test))
print("Mean squared error:",error)
'''

plt.figure()
plt.plot(range(len(cost_train)), np.reshape(cost_train, [len(cost_train), 1]), label = "Train Cost")
plt.plot(range(len(cost_test)), np.reshape(cost_test, [len(cost_test), 1]), label = "Test Cost")
plt.title("Cost/loss per iteration")
plt.xlabel("Number of iterations")
plt.ylabel("Cost/Loss")
plt.legend()
plt.show()


#error = cost_function(optimal_b, optimal_w, np.asmatrix(x_test), np.asmatrix(y_test))
#print("Mean squared error: %.2f" % error)


```

```

Coefficient: [[-0.41144214]
 [-1.32743691]
 [-1.6211677 ]
 [ 1.4974427 ]
 [ 0.34429883]
 [ 4.27231502]
 [-0.16421896]
 [-1.517363 ]
 [ 1.05404979]
 [-0.82986638]
 [ 0.22719132]
 [-0.03506302]
 [-3.44585941]]
y_intercept: [[21.27676176]]

```





Compare Custom SGD with Sklearn SGD

In [23]:

```
# Sklearn SGD
# The mean squared error
print("Mean squared error: %.2f" % mean_squared_error(Y_test, Y_pred))
# Explained variance score: 1 is perfect prediction
print("Variance score: %.2f" % r2_score(Y_test, Y_pred))
# The mean absolute error
print("Mean Absolute Error: %.2f" % mean_absolute_error(Y_test, Y_pred))
```

Mean squared error: 28.53
 Variance score: 0.70
 Mean Absolute Error: 3.46

In [24]:

```
# Implemented SGD
# The mean squared error
error = cost_function(optimal_b, optimal_w, np.asmatrix(x_test), np.asmatrix(y_test))
print("Mean squared error: %.2f" % (error))
# Explained variance score : 1 is perfect prediction
r_squared = r_sq_score(optimal_b, optimal_w, np.asmatrix(x_test), np.asmatrix(y_test))
print("Variance score: %.2f" % r_squared)
absolute_error = absolute_cost_function(optimal_b, optimal_w, np.asmatrix(x_test),
np.asmatrix(y_test))
print("Mean Absolute Error: %.2f" % absolute_error)
```



Mean squared error: 38.63
 Variance score: 0.36
 Mean Absolute Error: 1.77

In [25]:

```
# Scatter plot of test vs predicted
# sklearn SGD
plt.figure(1)
plt.subplot(211)
plt.scatter(Y_test, Y_pred)
plt.xlabel("Prices: $Y_i$")
plt.ylabel("Predicted prices: $\hat{Y}_i$")
plt.title("Prices vs Predicted prices: Sklearn SGD")
plt.show()


# Implemented SGD
plt.subplot(212)
plt.scatter([y_test], [(np.dot(np.asmatrix(x_test), optimal_w) + optimal_b)])
plt.xlabel("Prices: $Y_i$")
plt.ylabel("Predicted prices: $\hat{Y}_i$")
plt.title("Prices vs Predicted prices: Implemented SGD")
plt.show()
```






In [26]:

```
# Distribution of error
delta_y_im = np.asmatrix(y_test) - (np.dot(np.asmatrix(x_test), optimal_w) + optimal_b)
delta_y_sk = Y_test - Y_pred
import seaborn as sns;
import numpy as np;
sns.set_style('whitegrid')
sns.kdeplot(np.asarray(delta_y_im)[0], label = "Implemented SGD", bw = 0.5)
sns.kdeplot(np.array(delta_y_sk), label = "Sklearn SGD", bw = 0.5)
plt.title("Distribution of error:  $y_i - \hat{y}_i$ ")
plt.xlabel("Error")
plt.ylabel("Density")
plt.legend()
plt.show()
```



In [27]:

```
# Distribution of predicted value
sns.set_style('whitegrid')
sns.kdeplot(np.array(np.dot(np.asmatrix(x_test), optimal_w) + optimal_b).T[0], label = "Implemented SGD")
sns.kdeplot(Y_pred, label = "Sklearn SGD")
plt.title("Distribution of prediction  $\hat{y}_i$ ")
plt.xlabel("predicted values")
plt.ylabel("Density")
plt.show()
```





In [28]:

```
from prettytable import PrettyTable
# MSE = mean squared error
# MAE =mean absolute error
x=PrettyTable()#np.asmatrix(x_test),
x.field_names=['Model','Weight Vector','MSE','MAE', 'Variance Score']
x.add_row(['sklearn',sklearn_w,mean_squared_error(Y_test,
clf_.predict(X_test)),mean_absolute_error(Y_test, clf_.predict(X_test)),r2_score(Y_test, Y_pred)])
x.add_row(['custom',optimal_w,error,absolute_error,r_squared])
print(x)
```

Model	Weight Vector	MSE
MAE	Variance Score	
sklearn	[-1.19998397 0.70823758 -0.47154859 0.24990282 -1.25091027 2.89665502 28.545951396225185 3.450685306695088 0.6956551656111603	
	[-0.38369813 -2.64985472 1.78730939 -1.05960085 -2.08913235 1.03175028	
	[-3.31171865]	
custom	[[-0.41144214]	
	[38.62996824]] [[1.77117924]] [[0.3646231]]	
	[-1.32743691]	
	[-1.6211677]	
	[1.4974427]	
	[0.34429883]	
	[4.27231502]	
	[-0.16421896]	
	[-1.517363]	
	[1.05404979]	
	[-0.82986638]	
	[0.22719132]	
	[-0.03506302]	
	[-3.44585941]]	

In [29]:

```
sklearn_pred=clf_.predict(x_test)
implemented_pred=(np.dot(np.asmatrix(x_test), optimal_w) + optimal_b)
x=PrettyTable()
x.field_names=['SKLearn SGD predicted value','Implemented SGD predicted value']
for itr in range(15):
    x.add_row([sklearn_pred[itr],implemented_pred[itr]])
print(x)
```

SKLearn SGD predicted value	Implemented SGD predicted value
11.010976872064473	[[9.34267897]]
28.13265575430431	[[21.81915391]]
32.610429206840855	[[27.9704308411]]

	32.010429200040000		[[27.97040004]]	
	19.47265691695546		[[22.43740506]]	
	26.99547481859689		[[20.51530473]]	
	18.17885314254281		[[15.23322088]]	
	6.450183867637406		[[10.74021524]]	
	25.429866825378358		[[23.82642385]]	
	21.60484164577307		[[19.48934147]]	
	24.084364627932437		[[21.50028139]]	
	6.151923708168887		[[7.03308681]]	
	27.77673286644099		[[21.02386188]]	
	10.057460526020344		[[7.98370263]]	
	15.644363660603457		[[16.73402384]]	
	23.502153086425825		[[21.90299932]]	
+-----+-----+				